

Bloom filters are a variant on hash tables. They are called Bloom filters as they were developed by Burton Bloom in 1970. They are more space efficient than a hash table with the tradeoff there will be some false positive lookups (can be mitigated)

## SUPPORTED OPERATIONS

Fast lookups and inserts.

## COMPARISON TO HASH TABLES

PRO: More space efficient. This is because a Bloom filter only stores whether or not an object has been seen before, not the object or a reference to the object (pointer).

CONS:

1. Can't store an associated object. This is the tradeoff made for space efficiency.
2. No deletions are supported. Rather than perform a deletion which can lead to a false negative a new Bloom filter would be initialized and constructed.
3. Small false positive probability - The lookup says the item exists or has been inserted even though it has not been.

## APPLICATIONS

Useful for instances where space is at a premium or false positives will not adversely affect the program.

Original use: Spell checker. In the 1970s memory was at a premium and the English language is large

Canonical: Maintaining a list of forbidden passwords. Very quick to tell a user of an insufficiently secure password and the damage from a false positive is extremely minimal; the user just has to create another password.

Modern: Network routers. Not a whole lot of storage space and a router has to process a lot of information quickly: lookup of spurious IP addresses, maintain statistics to identify a denial of service attack, keep track of the contents of a cache to prevent slow disk lookups, etc.

## IMPLEMENTATION DETAILS

1. Array of  $n$  bits, bit = 0 or 1 (False or true, respectively)
2. Hash functions

## ARRAY STRUCTURE

The array consists of  $n$  bits.

$$n = \left( \begin{array}{l} \text{\# of bits chosen to} \\ \text{represent an object} \end{array} \right) \left( \begin{array}{l} \text{\# unique elements} \\ \text{in dataset } S \end{array} \right) \leftarrow \begin{array}{l} \text{cardinality of } S \\ |S| \end{array}$$

We can also adjust the size of the array to work within memory constraints.

$$\begin{array}{l} \text{\# of bits needed to} \\ \text{represent an object} \end{array} = n/|S|$$

While we can tune the size of the filter by varying  $n$  you get fewer false positives with decently sized  $n$ , for example  $n=8$  or higher.

## HASH FUNCTIONS

In a bloom filter there must be  $k$  hash functions ( $h_1, \dots, h_k$ ) with  $k$  being a small constant. The optimal value for  $k$  can be found with:

$$k = \left( \frac{\text{\# of bits in the array}}{\text{\# of inserted elements}} \right) \ln 2$$

## INSERT

To insert an element  $x$  into the Bloom filter  $A$ :

For  $i = 1, \dots, k$   
Set  $A[h_i(x)] = 1$  (Regardless if the value is 0 or 1 in  $A$ )

## LOOKUP

To lookup an element  $x$  in Bloom filter  $A$ :

Return True if  $A[h_i(x)] = 1$  for every  $i = 1, 2, \dots, k$

## IMPLEMENTATION DETAILS (cont)

Error Rate,  $\epsilon$ 

The third design factor is the design of a Bloom filter is knowing your acceptable error rate (rate of false positives) if there is a required constraint on the error rate it can have an effect on the size of the filter.

$$\epsilon \approx (1 - e^{-k\phi/n})^k$$

$k$  = number of hash functions

$\phi$  = number of objects inserted into the filter

$n$  = number of bits in (size of) the filter

As such, knowing the acceptable upper bound on  $\epsilon$  will help drive the parameters needed to correctly implement the Bloom filter.

$$n = \frac{-\phi \ln \epsilon}{(\ln 2)^2}$$

$$k = \frac{n}{\phi} \ln 2$$

3-0235 — 50 SHEETS — 5 SQUARES  
3-0236 — 100 SHEETS — 5 SQUARES  
3-0237 — 200 SHEETS — 5 SQUARES  
3-0137 — 200 SHEETS — FILLER

COMET